

---

# Resources & Resource Access Control

- ❑ Thus far...
  - ❑ We've assumed that tasks are independent of each other!
  - ❑ What if they share resources (other than processor itself)?
    - ❑ **Examples of resources**
      - ❑ Shared variables/data
      - ❑ Photo copier, printer, etc.
      - ❑ Remote server
  - ❑ So, what's different if tasks share resources
    - ❑ **Mutual exclusion** needs to be guaranteed
    - ❑ Scheduling needs to consider this aspect!

---

## What is “mutual exclusion”?

- ❑ Assurance that only one task accesses resource at a time
- ❑ Simple example scenario: consider tasks  $T_1$  and  $T_2$ 
  - ❑  $T_1$  (higher-priority task)
    - ❑ Withdraws money (say \$50) from bank account
  - ❑  $T_2$  (lower-priority task)
    - ❑ Deposits money (say \$50) into same bank account

---

# Mutual exclusion example

Account balance

100

Task  $T_1$ :

Load balance

Subtract withdrawal amount

Store updated balance

Task  $T_2$ :

Load balance

Add deposit amount

Store updated balance

---

---

# Mutual exclusion example

Account balance

100

Task  $T_1$ :

An empty rectangular box, likely representing a lock or semaphore used for mutual exclusion.

Load balance

Subtract withdrawal amount

Store updated balance

Task  $T_2$ :

100

Load balance

Add deposit amount

Store updated balance

---

---

# Mutual exclusion example

Account balance

100

Task  $T_1$ :

An empty rectangular box, likely representing a lock or semaphore used for mutual exclusion.

Load balance

Subtract withdrawal amount

Store updated balance

Task  $T_2$ :

150

Load balance

Add deposit amount

Store updated balance

---

---

# Mutual exclusion example

Account balance

100

Task  $T_1$ :

100

Load balance

Subtract withdrawal amount

Store updated balance

Task  $T_2$ :

150

Load balance

Add deposit amount

Store updated balance

---

---

# Mutual exclusion example

Account balance

100

Task  $T_1$ :

50

Load balance

Subtract withdrawal amount

Store updated balance

Task  $T_2$ :

150

Load balance

Add deposit amount

Store updated balance

---

---

# Mutual exclusion example

Account balance

50

Task  $T_1$ :

50

Load balance

Subtract withdrawal amount

Store updated balance

Task  $T_2$ :

150

Load balance

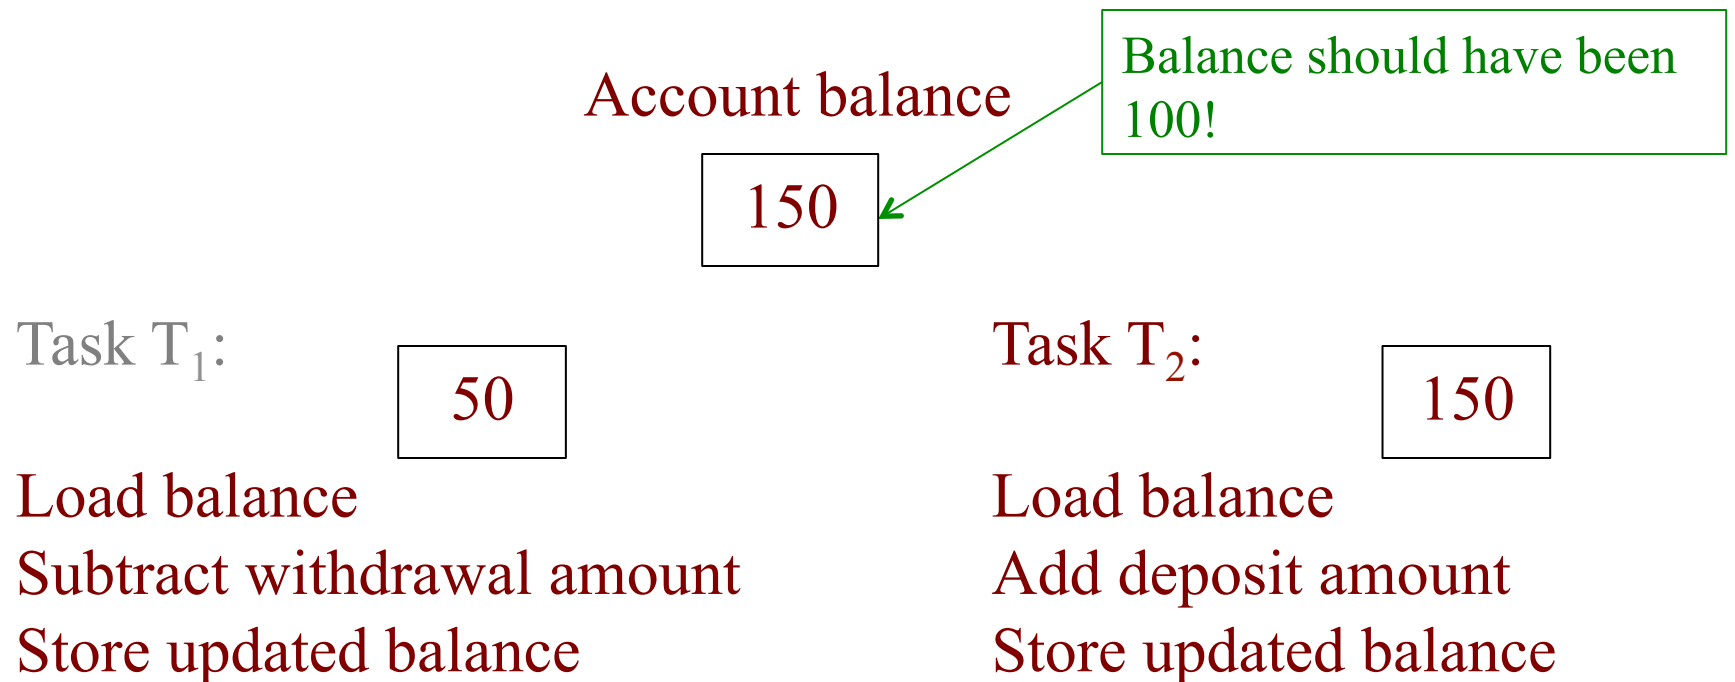
Add deposit amount

Store updated balance

---



# Mutual exclusion example



Conclusion: there should ***no interleaving*** of execution when tasks are accessing shared resources!!!

---

## So, to enforce mutual exclusion...

- Tasks must ***lock*** shared resources before using them
  - No other task must be granted resource until it is *unlocked*!

---

## Shared Resources

- Add to model: set of  $\rho$  serially **reusable resources**  $R_1, R_2, \dots, R_\rho$ , where there are  $v_i$  units of resource  $R_i$ .
  - **Examples of resources:**
    - Binary semaphore, for which there is one unit.
    - Counting semaphore, for which there may be many units.
    - Reader/writer locks.
    - Printer.
    - Remote server.

---

# Locks

- If  $n$  units of resource  $R$  are required
  - Lock request  $\rightarrow L(R, n)$
  - Corresponding unlock request  $\rightarrow U(R, n)$
- A matching lock/unlock pair is a **critical section**
- Critical section denoted by  $[R, n; e]$ 
  - $R \rightarrow$  resource requested
  - $n \rightarrow$  number of units
  - $e \rightarrow$  execution time of critical section

---

## Locks (Continued)

- Locks can be **nested**
- Notation:
  - $[R_1; 14 [R_4, 3; 9 [R_5, 4; 3]]]$
- Mostly interested in **outermost critical sections**
- **Note:** Assume there is only one kind of lock request
  - E.g., cannot distinguish reader locks/writer locks

---

# Conflicts and Contentions


## ■ Resource **conflict**

- Two jobs requesting common resources have conflict
- If we had reader/writer locks, then the notion of a “conflict” would be a little more complicated

## ■ Resource **contention**

- Manifestation of conflict during system operation
- One job requests for resource held by another job
- Requesting job must wait...it is said to be **blocked**  
→ need to calculate *upper bound on waiting/blocking time*

## Example of resource contention



 = access of single-unit resource R

---

## Dealing with resource contentions

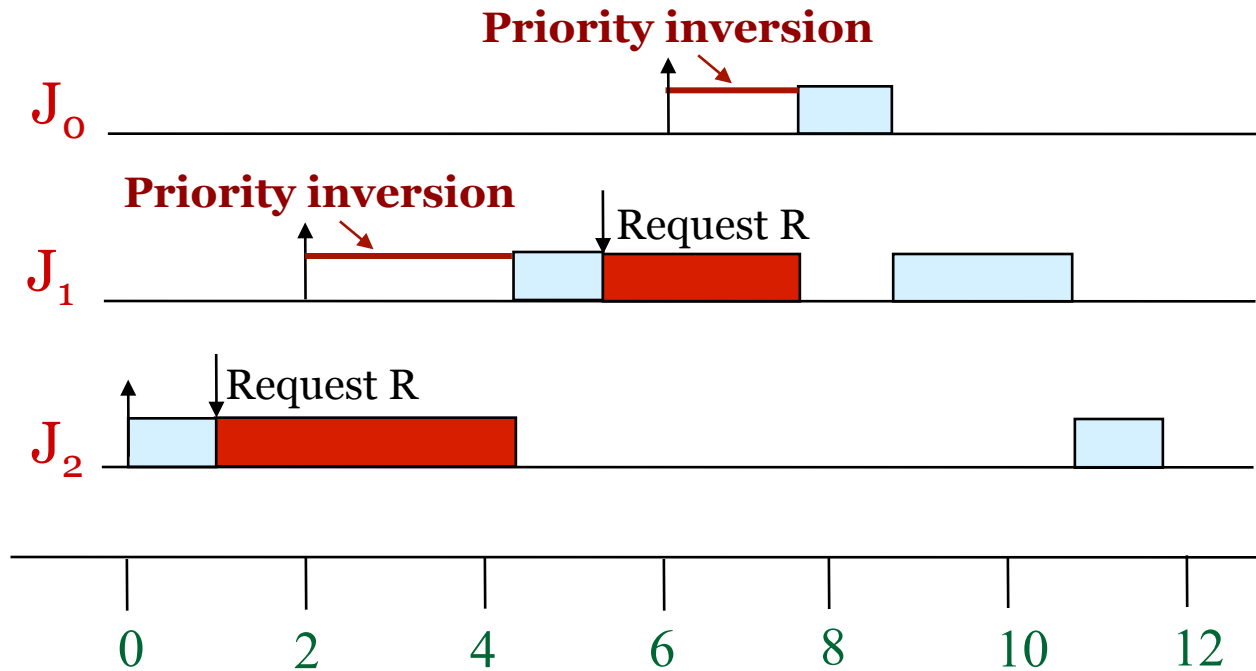
- Execute all critical sections **non-preemptively!**
  - If tasks are indexed by priority, **blocking term** of  $T_i$  is
$$\max_{i+1 \leq k \leq n} e_k$$
    - $e_k$  is the execution cost of the longest critical section of  $T_k$
  - **Note:** Critical section → **outermost critical section** unless otherwise specified



# Example




## Let's add another job into the mix...



What are the **disadvantages** here?

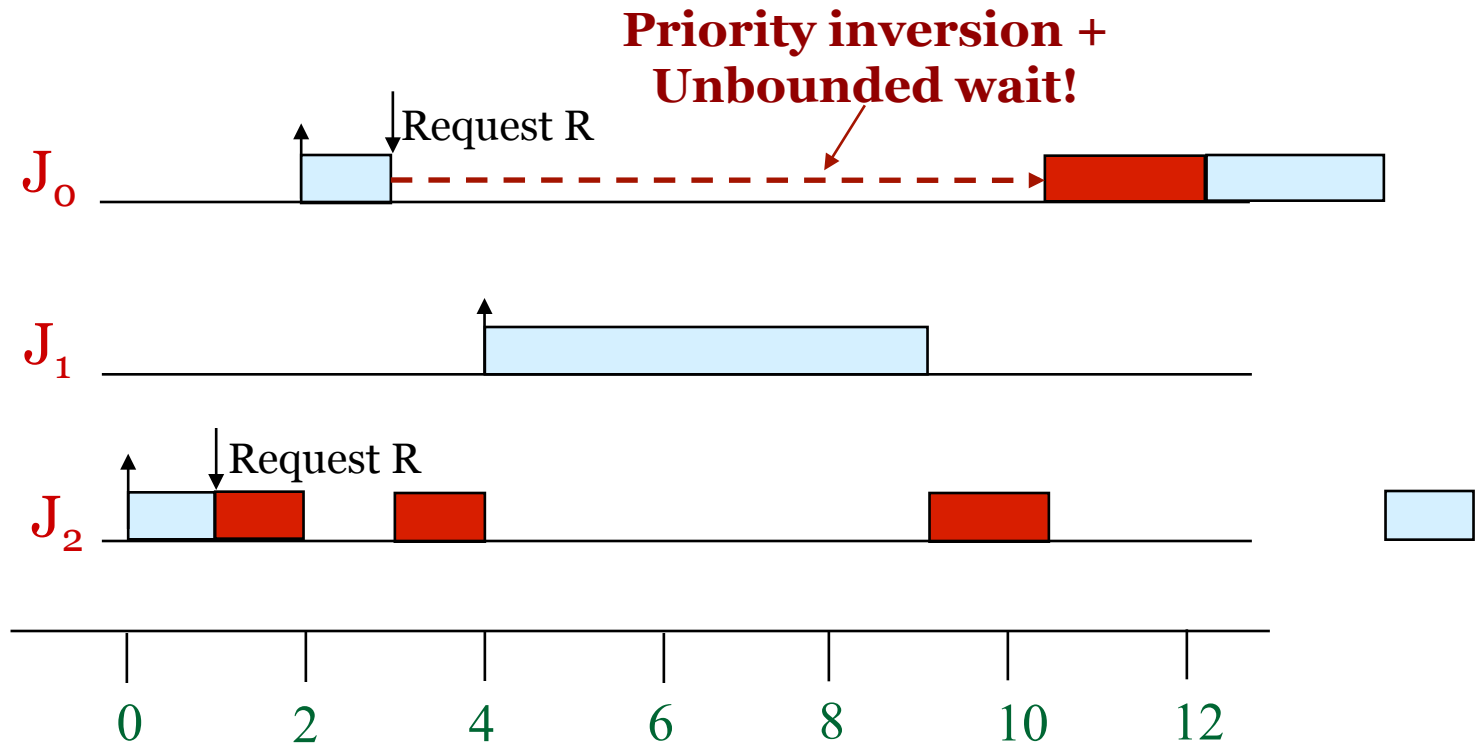
## Making only resources non-preemptive...

- Allow jobs with resources to be preempted, handle contention dynamically
  - i.e., start higher-priority jobs as they arrive...handle contention for R if and when it happens



Have we solved all issues?

## No, not quite! Consider another scenario...




The problem here is the medium priority job  $J_1$ !!!

---

## The problem, once again...

- Problems in previous example
  - Priority inversion
    - This can't be eliminated altogether when sharing resources
  - Unbounded wait time for  $J_0$ 
    - Caused due to medium-priority task  $J_1$
    - Fact that  $J_0$  is waiting is not considered by scheduler
      - Scheduler only knows about priorities...
    - Solution?
      - Fact that  $J_0$  is waiting needs to be reflected in priorities!
      - Make  $J_2$  “inherit”  $J_0$ 's priority while using resource
      - “Priority Inheritance Protocol”

## Example revisited...



---

# Priority Inheritance Protocol Definition

Each job  $J$  has **assigned priority** and **current priority**. At its release time, current priority of every job is equal to its assigned priority.

1. **Scheduling Rule:** Ready jobs scheduled **preemptively** in a **priority-driven** manner according to their *current* priorities.
2. **Allocation Rule:** When job  $J$  requests for resource  $R$  at time  $t$ 
  - a. If  $R$  is free,  $R$  is allocated to  $J$  until  $J$  releases it
  - b. If  $R$  is not free, the request is denied and  $J$  is *blocked*
3. **Priority-Inheritance Rule**
  - a. When requesting job  $J$  becomes *blocked*, job  $J_l$  that blocks  $J$  *inherits* the *current priority* of  $J$ .
  - b. Job  $J_l$  executes at its inherited priority until it releases  $R$  (or until it inherits an even higher priority)
  - c. Upon release of  $R$ , priority of  $J_l$  returns to priority it had when it acquired resource  $R$

---


## Formal terminology...

- ❑ Multiple serially **reusable resources**  $R_1, R_2, \dots, R_p$ , where there are  $v_i$  units of resource  $R_i$
- ❑ If  $n$  units of resource  $R$  are required
  - ❑ Lock request  $\rightarrow \mathbf{L(R, n)}$
  - ❑ Corresponding unlock request  $\rightarrow \mathbf{U(R, n)}$
- ❑ A matching lock/unlock pair is a **critical section**
- ❑ Critical section denoted by  $\mathbf{[R, n; e]}$ 
  - ❑  $R \rightarrow$  resource requested
  - ❑  $n \rightarrow$  number of units
  - ❑  $e \rightarrow$  execution time of critical section



## Nested resource usage

- ❑ Locks can be **nested**
- ❑ Notation:
  - ❑  $[R_1; 14 [R_4, 3; 9 [R_5, 4; 3]]]$




By allowing transitive priority inheritance,  
the protocol we have discussed extends  
to nested resource usage as well

## A more complicated example

Job	Release time	Execution time	Critical sections
J1	7	3	[Green; 1] (from 1)
J2	5	3	[Red;1] (from 1)
J3	4	2	
J4	2	6	[Green; 4 [Red; 1.5]] (from 1, 2)
J5	0	6	[Red; 4] (from 1)

## Here's the schedule according to PIP



---

## Properties of PIP

- ❑ We have two kinds of blocking with the PIP: **direct blocking** and **inheritance blocking**
  - ❑  $J_2$  directly blocked by  $J_5$  during interval (6,10]
  - ❑  $J_2$  inheritance blocked by  $J_4$  during interval (10,13]
  - ❑  $J_3$  inheritance blocked by  $J_5$  during interval (6,7]
- ❑ Jobs can **transitively** block each other
  - ❑ @ 8,  $J_5$  blocks  $J_4$  and  $J_4$  blocks  $J_1$

---

## PIP Schedulability Analysis

**Lemma 1:** Job  $J_k$  waits for at most one completion of a critical section of job that blocks it (directly or indirectly), **regardless of how many times  $J_k$  is using each lock.**

**Proof sketch:** since  $k$  has higher priority, once  $J_i$  gets out,  $J_k$  will not let it begin something else.

**Lemma 2:** If job  $k$  uses lock  $s$ , then it can be blocked for the duration of at most one critical section guarded by  $s$ , regardless how many times  $s$  is used by how many jobs.

**Proof sketch:**  $k$  has higher priority (by definition) than all jobs that try to block it. Hence, the first time  $k$  tries to get a lock  $s$ , it will be blocked for at most the time it takes for the lower-priority job inside to get out. The following times, no lower-priority job may have gotten inside a section guarded by  $s$ .

## A slight modification...

Job	Release time	Execution time	Critical sections
J1	7	3	[ <b>Green</b> ; 1] (from 1)
J2	5	3	[ <b>Red</b> ;1] (from 1)
J3	4	2	
J4	2	6	[ <b>Green</b> ; 4 [ <b>Red</b> ; 1.5]] (from 1, 2)
J5	0	6	[ <b>Red</b> ; 4 [ <b>Green</b> ; 1]] (from 1, 2)

---

## Locks and Deadlocks

J<sub>5</sub>:  
do {

lock (Red)

lock (Green)

use (Red, Green)

unlock (Green)

unlock (Red)

}

J<sub>4</sub>:  
do {

lock (Green)

lock (Red)

use (Red, Green)

unlock (Red)

unlock (Green)

}

Possible result: J<sub>5</sub> gets **Red**, J<sub>4</sub> gets **Green** → deadlock!

How can we fix the problem?

---

## Need Partial/Total Order of Locks

- ❑ Choose a total order of all locks:  $L_1 < L_2 < \dots < L_n$ 
  - ❑ Or a partial order for subsets of locks used by same tasks, e.g.,  $L_1 < L_3 < L_4$  and  $L_2 < L_5$
- ❑ When you need both  $L_i$  and  $L_j$ , get that lock first that comes first in the order



---


## Is this the best protocol for high-priority tasks?

- ❑ Assume we know what tasks request what resources
  - ❑ Can this info. be used to the advantage of high-prio tasks?
  - ❑ If yes, how?

## Let's revisit our PIP example...

Job	Release time	Execution time	Critical sections
J1	7	3	[Green; 1] (from 1)
J2	5	3	[Red;1] (from 1)
J3	4	2	
J4	2	6	[Green; 4 [Red; 1.5]] (from 1, 2)
J5	0	6	[Red; 4] (from 1)

## Here's what the timeline looked like




---

## Is this the best we can do for high-prio tasks?

- ❑ Assume we know when and what resources tasks request
  - ❑ Can this info. be used to reduce blocking of high-prio tasks?
  - ❑ If yes, how?
  - ❑ Basic idea
    - ❑ Identify highest priority task using every resource
      - ❑ Resource ceiling
    - ❑ Keep track of resources currently in use
      - ❑ System ceiling
    - ❑ Prevent locking of resource by task that has lower priority than system ceiling

# How does the schedule change?




## Yet another “complicated” example

Job	Release time	Execution time	Critical sections
J1	7	3	[DarkGreen; 1] (from 1), [LightGreen; 1] (from 2)
J2	5	3	[DarkGreen; 1] (from 1)
J3	4	2	
J4	2	6	[LightGreen; 4 [Red; 1]] (from 1, 2)
J5	0	6	[Red; 4 [DarkGreen; 3]] (from 1, 2)



$$DG \rightarrow \pi, \quad LG \rightarrow \pi, \quad R \rightarrow \pi_4$$





---

# The Priority Ceiling Protocol (Sha, Rajkumar, Lehoczky)

## □ Two key assumptions:

- The assigned priorities of all jobs are fixed (as before).
- The resources required by all jobs are known *a priori* before the execution of any job begins

**Definition:** The **priority ceiling** of any resource  $R$  is the highest priority of all the jobs that require  $R$ , and is denoted by  $\Pi(R)$

**Definition:** The **current priority ceiling**  $\Pi'(t)$  of the system is equal to the highest priority ceiling of the resources currently in use, or  $\Omega$  if no resources are currently in use ( $\Omega$  is a priority lower than any real priority)



---

# PCP Definition

## 1. Scheduling Rule:

(a) At release time  $t$ , current priority  $\pi(t)$  of every job  $J$  = assigned priority. Job remains at this priority except under conditions of rule 3.

(b) Jobs  $J$  scheduled preemptively + strict priority at its current priority  $\pi(t)$ .

## 2. Allocation Rule: Whenever a job $J$ requests a resource $R$ at time $t$ , either:

(a)  $R$  is held by another job --  $J$ 's request fails +  $J$  becomes blocked -- or

(b)  $R$  is free.

(i) If  $J$ 's prio  $\pi(t) >$  current prio ceiling  $\Pi'(t)$ ,  $R$  is allocated to  $J$ .

(ii) If  $J$ 's prio  $\pi(t) \leq$  ceiling  $\Pi'(t)$ ,  $R$  is allocated to  $J$  only if  $J$  is the job holding the resource(s) whose priority ceiling =  $\Pi'(t)$ ; otherwise,  $J$ 's request is denied and  $J$  becomes blocked.

**3. Priority-Inheritance Rule:** When  $J$  becomes blocked, the job  $J_l$  that blocks  $J$  inherits the current priority  $\pi(t)$  of  $J$ .  $J_l$  executes at its inherited priority until it releases every resource whose priority ceiling is  $\geq \pi(t)$  (or until it inherits an even higher priority); at that time, the priority of  $J_l$  returns to the priority  $\pi_l(t')$  it had at the time  $t'$  when it was granted the resources.